

SOLOMON SMITH

Montclair, CA | (530) 513-1936 | solomonsmithdev@gmail.com
linkedin.com/in/solomonsmithdev | github.com/SolomonSmith-dev | solomonsmith.dev
AI/Backend Engineer | B.S. Computer Science, Dec 2026 | Available January 2027

EDUCATION

California State University, San Bernardino

B.S. Computer Science

Dec 2026

San Bernardino, CA

Relevant coursework: Machine Learning, Algorithms and Data Structures, Database Systems, Operating Systems

EXPERIENCE

Founding Engineer (2nd engineer)

Jul 2026 to present

Recura AI / CourtRules (courtrules.app), legal-data platform covering filing rules for 1,653 judges

Remote

- Shipped retrieval-grounded Q&A on judge pages, which carry 57% of site pageviews: fetches filed court orders from object storage, extracts per-page text, grounds the model in it, and streams answers as NDJSON (Next.js, Supabase).
- Designed a fail-closed citation verifier: a claim renders only if its quote appears verbatim in the source. Its tests caught a bypass where an invented quote passed on a plausible source index, fixed before release.
- Built query analytics that classify questions into a 14-type rule taxonomy, plus a Postgres view ranking unanswered questions to separate data gaps from retrieval failures (service-write-only RLS).
- Added iCal and CSV court-calendar exports (RFC 5545, RFC 4180) with a formula-injection guard (CWE-1236).
- Wrote the first README and automated the shared-package build so a fresh clone runs with no manual steps.

Full-Stack Developer (Contract)

Jun 2025 to Sep 2025

RideSplits

San Bernardino, CA

- Migrated auth from custom JWT to Firebase Auth for rider and driver roles with role-based Firestore rules.
- Built a multi-step ID-verification upload flow with client-side file-type and size validation.
- Documented the role-based Firestore and Storage security rules for the team.

PROJECTS

ARDA, multi-agent service | *LangGraph, LlamaIndex, Redis, MCP*

github.com/SolomonSmith-dev/arda

- Built a LangGraph orchestrator on native Anthropic tool_use with a Redis-backed executor, LlamaIndex RAG, and an MCP server behind 17 documented FastAPI routes.
- Designed a mock LLM and embedder layer so the 425-test suite runs offline with no API keys in 13 seconds; CI gates on ruff, mypy, and pytest.
- Operate it 24/7 on a self-hosted Debian server (PM2, Tailscale); startup fails closed on a missing API key.

SOC Triage AI, RAG alert triage | *Claude, Streamlit*

github.com/SolomonSmith-dev/soc-triage-ai

- Built a six-stage triage pipeline: regex observable extraction, retrieval over MITRE ATT&CK, context-constrained Claude generation, and strict JSON schema validation.
- Added a similarity guardrail that refuses out-of-scope alerts before any generation runs.
- Raised reliability-harness pass rate from 3 of 7 to 7 of 7 cases by reworking corpus chunking; shipped a Streamlit analyst UI with evidence panels and override tracking.

PhishGuard | *LightGBM, ONNX Runtime, FastAPI*

github.com/SolomonSmith-dev/phishguard

- Found and fixed label-polarity and distribution leakage in the PhiUSIIL dataset (every legitimate URL began with https://www). Documented the methodology in LIMITATIONS.md and added leakage tests to pre-commit.
- Trained a LightGBM URL classifier to 0.9943 test AUC at a 1.54% false-positive rate on Tranco top-5,000 domains; served via FastAPI with ONNX Runtime.

SKILLS

Backend/Infra: FastAPI, Next.js, Node.js, PostgreSQL/Supabase, Redis, Firebase, Docker, GitHub Actions, Linux

Languages: Python, TypeScript, SQL, C++, Bash

AI/ML: LightGBM, scikit-learn, ONNX Runtime, LangGraph, LlamaIndex, Anthropic SDK, sentence-transformers

PRIOR CAREER

Professional kitchens, 10+ years: Dishwasher promoted to Head Chef